



SESC

SUPER ESCALAR SIMULATOR

Joe Renau, Basilio Fraguera, and Liu Wei



University of Illinois at Urbana-Champaign

<http://sesc.sourceforge.net>

March 2002



TM



JOE VS. THE SIMULATORS

94-95 SIMUS: An R2000 trace-driven sim (14KLOC)



JOE VS. THE SIMULATORS

- 94-95** SIMUS: An R2000 trace-driven sim (14KLOC)
- 98-01** MINT+SMT: An R3000 execution-driven sim
- MINT: Real MIPS R3000 execution (26KLOC)
 - SMT: Timing simulator (19KLOC)
 - FlexRAM: A sample memory backend (3.5KLOC)



JOE VS. THE SIMULATORS

- 94-95** SIMUS: An R2000 trace-driven sim (14KLOC)
- 98-01** MINT+SMT: An R3000 execution-driven sim
- MINT: Real MIPS R3000 execution (26KLOC)
 - SMT: Timing simulator (19KLOC)
 - FlexRAM: A sample memory backend (3.5KLOC)
- 01-02** SESC: An R3000 execution-driven sim
- MINT: Real MIPS R3000 execution (26KLOC)
 - SESC: Timing simulator in C++ (22KLOC)



JOE VS. THE SIMULATORS

- 94-95** SIMUS: An R2000 trace-driven sim (14KLOC)
- 98-01** MINT+SMT: An R3000 execution-driven sim
- MINT: Real MIPS R3000 execution (26KLOC)
 - SMT: Timing simulator (19KLOC)
 - FlexRAM: A sample memory backend (3.5KLOC)
- 01-02** SESC: An R3000 execution-driven sim
- MINT: Real MIPS R3000 execution (26KLOC)
 - SESC: Timing simulator in C++ (22KLOC)
- SimpleScalar 4.0 (55KLOC)





WHY ANOTHER ONE?

- Model more details in the out-of-order core
- The need for speed
- I do not fully understand smt (>3 years)





WHY ANOTHER ONE?

- Model more details in the out-of-order core
- The need for speed
- I do not fully understand smt (>3 years)
- To make the new simulator understandable:
 - Clean C++ design
 - Document in source
 - Remove the weird MINT scheduling
 - Common powerful configuration format
 - Facilitate code sharing between projects



OBJECTIVE

- Parameters:
 - Multiprocessor with Release Consistency
 - Multiple types of processors simulatenously
 - SMTs
 - Add some “kind-of OS”: TLB, process. . .



OBJECTIVE

- Parameters:
 - Multiprocessor with Release Consistency
 - Multiple types of processors simulatenously
 - SMTs
 - Add some “kind-of OS”: TLB, process...
- Results:
 - Separate stats generation and processing:
 - The simulator outputs close to raw data
 - Perl scripts process the data
 - Profile applications at inst. level





MOTTOS



- Build the fastest possible simulator
 - Every addition is evaluated
 - Exploit the most common pattern
 - Checks only in debug mode



MOTTOS

- Build the fastest possible simulator
 - Every addition is evaluated
 - Exploit the most common pattern
 - Checks only in debug mode
- Make the code easy to understand and extend
 - Be ashamed of bad code
 - Uniform coding style



MOTTOS

- Build the fastest possible simulator
 - Every addition is evaluated
 - Exploit the most common pattern
 - Checks only in debug mode
- Make the code easy to understand and extend
 - Be ashamed of bad code
 - Uniform coding style
- Provide as many configurations as possible
 - Add all the reasonable parameters



CODING STYLE

- Indentation (2 tabstop)
- Braces (Linux Kernel style)

```
if( x ) {  
    }  
}
```
- Assertions: Use nanassert wherever possible
- Comments, do not overcomment. Use doxygen
- Naming convention
- Code replication (rule of 3)



MAIN DIRECTORIES

src main core of the simulator

backend memory sub-system (based on FlexRAM)

os OS support (TLB)

libapp library necessary for applications (ala libsim.a)

doc documentation files and doxygen output


benchs sample bench and configuration

scripts for processing the statistics generated

FlexRAM the FlexRAM architecture (not in sourceforge)


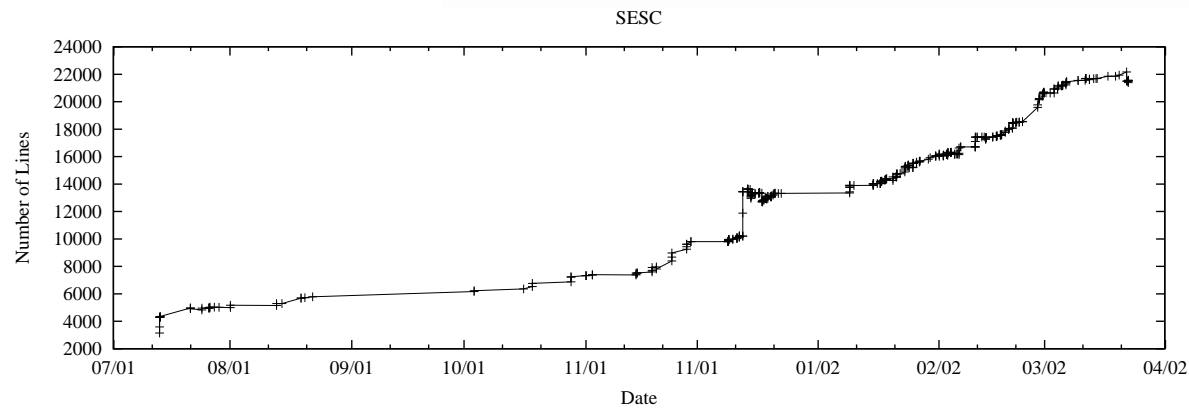


CVS WORK





CVS WORK



- Thread support:
spawn, wait, self, suspend, resume, yield, exit

- Thread support:
spawn, wait, self, suspend, resume, yield, exit
- Barriers and locks:
lock/unlock, barrier, psemaphore/vsemaphore, C4 Macros

- Thread support:
spawn, wait, self, suspend, resume, yield, exit
- Barriers and locks:
lock/unlock, barrier, psemaphore/vsemaphore, C4 Macros
- Events:
 - preevent: Notified when executed in front-end
 - postevent: Goes through the pipeline
 - memfence : Previous memory operations issued
 - acquire,release : Release Consistency
 - fast_sim_begin, fast_sim_end

MEMORY SUBSYSTEM

```
class DCache:public MemoryObj {
protected:
    void read() {
        MSG("0x%x:ld [0x%x]",req->getIaddr(),req->getVaddr());
        req->goUp(1);
    };

    void write() {
        MSG("time %lld",globalClock);
        req->goUp(1);
    };

    void memFence() { req->goUp(23); };
    void acquire() { req->goUp(10); };
    void release() { req->goUp(13); };
};
```



BOOTING SESC

```
int main(int argc, char **argv, char **envp) {  
  
    DummyMemorySystem *cm[NPROC];  
    MyProc *core[NPROC];  
  
    osSim = new OSSim(argc, argv, envp);  
  
    for(Pid_t i = 0; i < NPROC; i++) {  
        cm[i] = new DummyMemorySystem(i);  
        core[i] = new MyProc(cm[i], i);  
    }  
  
    osSim->boot();  
  
    for(int i = 0; i < NPROC; i++) {  
        delete core[i];  
        delete cm[i];  
    }  
    delete osSim;  
}
```





RUNNING SESC



- Compile a program for O32 in SGI (link libapp)





RUNNING SESC



- Compile a program for O32 in SGI (link libapp)
- Check configuration file (sesc.conf)



RUNNING SESC

- Compile a program for O32 in SGI (link libapp)
- Check configuration file (sesc.conf)
- `./mtst1 -h0x8000000 crafty <tt.in`



RUNNING SESC

- Compile a program for O32 in SGI (link libapp)
- Check configuration file (sesc.conf)
- `./mtst1 -h0x8000000 crafty <tt.in`
- Check output `sesc_crafty.xxxxxxx`



RUNNING SESC

- Compile a program for O32 in SGI (link libapp)
- Check configuration file (sesc.conf)
- `./mtst1 -h0x8000000 crafty <tt.in`
- Check output `sesc_crafty.xxxxxxx`
- Process it `../scripts/report.pl sesc_crafty.xxxxxxx`



SAMPLE CONFIGURATION

```
foo = 4
cpucore[0:31] = 'r10' # iacoma3
[r10]
fetchWidth = 4
issueWidth = $(foo)
branchDelay = $(foo)+2 # 6 minimum (deps enforced)
maxBranches = 4 # outstanding branches limit
maxIRequests= 2
iALUUnits = 2
iFPUnits = 2
iBJUnits = 1
LDSTQueue = 16 # Unified LD/ST buffer queue
winsize = 32
bb4Cycle = 1 # max BB fetch per cycle
bpred = 'r10kBPred'

<../benchs/branchModels.conf> # Include different branch models
```



REPORT FORMAT

- By default is sesc_bench.xxxxxxx
- Includes the input configuration file
- Homogenous way to report data (GStats.h)

KEY:field1=value1:field2=value2:...

- Example:

```
ProcessId(0):totalTime=358100288:waitTime=0:spawnTime=0
Proc(0):clockTicks=358100289
BPred_RAS(0):nHit=      13
BPred_RAS(0):nMiss=     0
cache(0):readHit=      6788685
```



REPORT OUTPUT

- Simulating an R10K at 180Mhz
- Crafty native execution 109msec

```
# Bench : ./mtst1 -h0x800000 ../benchs/crafty
# File  : sesc_crafty.EMomhe      :      Mon Mar 25 21:27:33 2002
      Sim. Speed      Exe Time      Sim Time (180MHz)
      1137.96 KIPS      17.88 secs      105.50 msec
Proc  Type      Total      RAS      BPred      BTB
   0  2bit      80.819% ( 99.91% of 10.05%) 84.02% 77.40%
File  IPC      Cycles      Busy      Control      Window      Struct      Memory
   0  1.07  18990598  26.79%  29.29%  18.45%  15.42%  0.20%  9.85%
```



PROFILING

Function	nCalls	call(clk)	IPC	Acum(%)	Time(%)	Ctrl(%)
SCVQScores_al	998	1161690.1	1.19	23.34%	23.34%	25.79%
chan_v_eval	1568385	578.9	0.43	41.62%	18.28%	72.29%
_doscan	903173	581.6	1.56	52.20%	10.58%	68.07%
skipto	22804	9285.8	1.05	56.46%	4.26%	78.42%
hmm_tied_read	47	4469093.4	1.83	60.69%	4.23%	62.13%
root_chan_v_	386348	488.0	0.60	64.48%	3.80%	63.28%
.....						



SIMULATION ACCURACY I

Simulating an R10K at 180Mhz (SGI Origin 200 IP27):

Bench	Native	Simulated	Error	SimTime
matrix	78ms	80ms	+2%	18s
crafty	109ms	105ms	-3%	18s
mcf	791ms	792ms	+0%	142s
mp3dec	643ms	619ms	-5%	160s
lat_mem	2405ms	2532ms	+5%	230s

SIMULATION ACCURACY II

Simulating an R4400 at 150Mhz (SGI IP22):

Bench	Native	Simulated	Error	SimTime
matrix	291ms	282ms	-4%	27s
crafty	265ms	271ms	+4%	43s
mcf	2438ms	2422ms	-1%	174s
mp3dec	2347ms	2521ms	+7%	185s
lat_mem	4005ms	4129ms	+3%	279s



SIMULATION SPEED

- Crafty simulation in a Pentium-III 1GHz

1137KIPS SESC simulating an R10K (with backend)



SIMULATION SPEED

- Crafty simulation in a Pentium-III 1GHz

1137KIPS SESC simulating an R10K (with backend)

178KIPS smt simulating an R10K (no backend)



SIMULATION SPEED

- Crafty simulation in a Pentium-III 1GHz

1137KIPS SESC simulating an R10K (with backend)

178KIPS smt simulating an R10K (no backend)

2608KIPS mint in-order



SIMULATION SPEED

- Crafty simulation in a Pentium-III 1GHz

1137KIPS SESC simulating an R10K (with backend)

178KIPS smt simulating an R10K (no backend)

2608KIPS mint in-order

≈150KIPS simplescalr sim-outorder





FULL SPECINT



- Assuming 1000KIPS we can run:
- 175.vpr Ref in $\approx 30h$



FULL SPECINT

- Assuming 1000KIPS we can run:
- 175.vpr Ref in $\approx 30h$
- 197.parser Ref in $\approx 90h$



FULL SPECINT

- Assuming 1000KIPS we can run:
- 175.vpr Ref in $\approx 30h$
- 197.parser Ref in $\approx 90h$
- Reduced Data sets (Minnesota) in $\approx 30min$



EVEN FASTER

- Parallelization:
 - Low ≈ 2 : 2 threads likely to be balanced
 - High ≈ 4 : Several threads in multiprocessors



EVEN FASTER

- Parallelization:
 - Low ≈ 2 : 2 threads likely to be balanced
 - High ≈ 4 : Several threads in multiprocessors
- Configuration Value Profile $\approx 20\% - 40\%$
 - Convert variables to constants
 - The compiler simplifies the code
 - Run-time verification

Close to 3MIPS in our dual athlon



EVEN FASTER

- Parallelization:
 - Low ≈ 2 : 2 threads likely to be balanced
 - High ≈ 4 : Several threads in multiprocessors
- Configuration Value Profile $\approx 20\% - 40\%$
 - Convert variables to constants
 - The compiler simplifies the code
 - Run-time verification

Close to 3MIPS in our dual athlon



DO YOU WANNA RUN IT?

- Add <http://sesc.sourceforge.net> to your bookmark
- Subscribe to sesc-news@lists.sourceforge.net
- Read all the files in doc directory
- See the ./src/main.cpp example
- Understand the callbacks (DInst.h is easy)
- Understand the configuration files (*.conf)



- More advance branch predictor: Multi-level, multi-cycle

- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)

- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)

TODO

- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them



- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them
- Type a hackers guide (started)

- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them
- Type a hackers guide (started)
- Energy model?



- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them
- Type a hackers guide (started)
- Energy model?
- Give more flexibility to profiling



- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them
- Type a hackers guide (started)
- Energy model?
- Give more flexibility to profiling
- Report sync time in multiprocessor

