



# SESC

## SUPER ESCALAR SIMULATOR

Joe Renau, Basilio Fraguera, and Liu Wei



University of Illinois at Urbana-Champaign

<http://sesc.sourceforge.net>

March 2002



TM



# OUTLINE



- Why Another One?
- Objective
- Mottos
- Coding Style
- Code Organization
- SESC Output
- Simulation Accuracy and Speed
- Todo





# WHY ANOTHER ONE?

- Model more details in the out-of-order core
- The need for speed
- After 3 years I still do not understand smt





# WHY ANOTHER ONE?

- Model more details in the out-of-order core
- The need for speed
- After 3 years I still do not understand smt
- To make the new simulator understandable:
  - Clean C++ design
  - Document in source
  - Remove the weird MINT scheduling
  - Common powerful configuration format
  - Facilitate code sharing between projects

# OBJECTIVE

- Parameters:
  - Multiprocessor with Release Consistency
  - Multiple types of processors concurrently
  - SMTs
  - Add some “kind-of OS”: TLB, process...



# OBJECTIVE

- Parameters:
  - Multiprocessor with Release Consistency
  - Multiple types of processors concurrently
  - SMTs
  - Add some “kind-of OS”: TLB, process...
- Results:
  - Separate stats generation and processing:
    - The simulator outputs close to raw data
    - Perl scripts process the data
  - Profile applications at inst. level



# MOTTOS

- Build the fastest possible simulator
  - Every addition is evaluated
  - Exploit the most common pattern
  - Checks only in debug mode



# MOTTOS

- Build the fastest possible simulator
  - Every addition is evaluated
  - Exploit the most common pattern
  - Checks only in debug mode
- Make the code easy to understand and extend
  - Be ashamed of bad code
  - Uniform coding style





# MOTTOS

- Build the fastest possible simulator
  - Every addition is evaluated
  - Exploit the most common pattern
  - Checks only in debug mode
- Make the code easy to understand and extend
  - Be ashamed of bad code
  - Uniform coding style
- Provide as many configurations as possible
  - Add all the reasonable parameters



# CODING STYLE

- Indentation (2 tabstop)
- Braces (Linux Kernel style)  

```
if( x ) {  
    }  
}
```
- Assertions: Use nanassert wherever possible
- Comments, do not over comment. Use doxygen
- Naming convention
- Code replication (rule of 3)



# MAIN DIRECTORIES

**src** main core of the simulator

**backend** memory sub-system (based on FlexRAM)

**os** OS support (TLB)

**libapp** library necessary for applications (ala libsim.a)

**doc** documentation files and doxygen output

**benchs** sample bench and configuration


**scripts** for processing the statistics generated

**FlexRAM** the FlexRAM architecture (not in sourceforge)



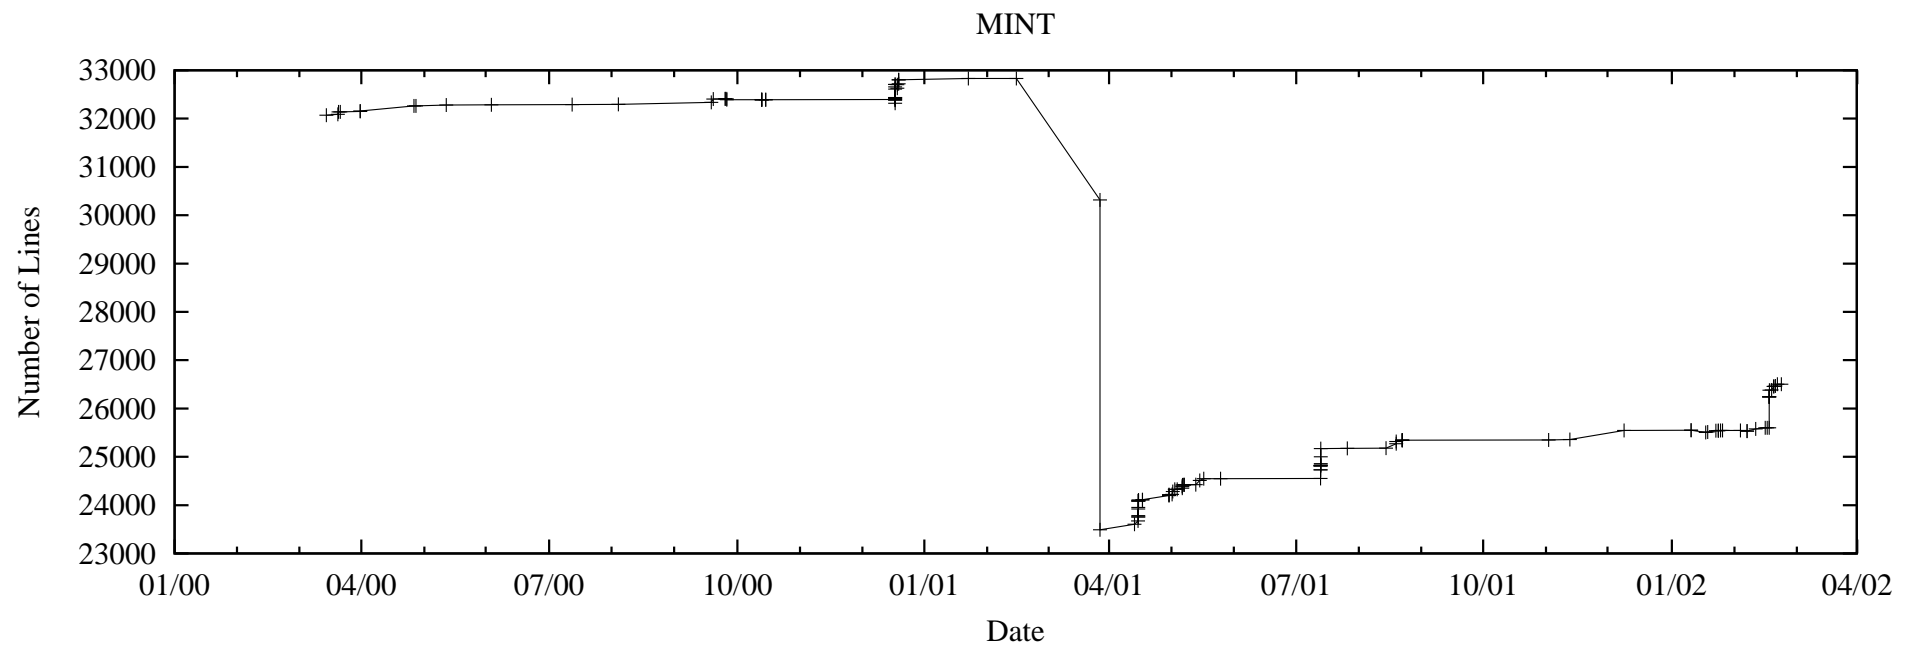



# CVS WORK





# CVS WORK



- Thread support:  
spawn, wait, self, suspend, resume, yield, exit

- Thread support:  
spawn, wait, self, suspend, resume, yield, exit
- Barriers and locks:  
lock/unlock, barrier, psemaphore/vsemaphore, C4 Macros

- Thread support:  
spawn, wait, self, suspend, resume, yield, exit
- Barriers and locks:  
lock/unlock, barrier, psemaphore/vsemaphore, C4 Macros
- Events:
  - preevent: Notified when executed in front-end
  - postevent: Goes through the pipeline
  - memfence : Previous memory operations issued
  - acquire,release : Release Consistency
  - fast\_sim\_begin, fast\_sim\_end



# MEMORY SUBSYSTEM

```
class DCache:public MemoryObj {
protected:
    void read() {
        MSG("0x%x:ld [0x%x]",req->getIaddr(),req->getVaddr());
        req->goUp(1);
    };

    void write() {
        MSG("time %lld",globalClock);
        req->goUp(1);
    };

    void memFence() { req->goUp(23); };
    void acquire() { req->goUp(10); };
    void release() { req->goUp(13); };
};
```



# BOOTING SESC

```
int main(int argc, char **argv, char **envp) {

    DummyMemorySystem *cm[NPROC];
    MyProc *core[NPROC];

    osSim = new OSSim(argc, argv, envp);

    for(Pid_t i = 0; i < NPROC; i ++) {
        cm[i] = new DummyMemorySystem(i);
        core[i] = new MyProc(cm[i], i);
    }

    osSim->boot();

    for(int i = 0; i < NPROC; i ++) {
        delete core[i];
        delete cm[i];
    }
    delete osSim;
}
```





# RUNNING SESC



- Compile a program for O32 in SGI (link libapp)



# RUNNING SESC

- Compile a program for O32 in SGI (link libapp)
- Check configuration file (sesc.conf)



# RUNNING SESC

- Compile a program for O32 in SGI (link libapp)
- Check configuration file (sesc.conf)
- `./mtst1 -h0x8000000 crafty <tt.in`



# RUNNING SESC

- Compile a program for O32 in SGI (link libapp)
- Check configuration file (sesc.conf)
- `./mtst1 -h0x8000000 crafty <tt.in`
- Check output `sesc_crafty.xxxxxxx`



# RUNNING SESC

- Compile a program for O32 in SGI (link libapp)
- Check configuration file (sesc.conf)
- `./mtst1 -h0x8000000 crafty <tt.in`
- Check output `sesc_crafty.xxxxxxx`
- Process it `../scripts/report.pl sesc_crafty.xxxxxxx`



# SAMPLE CONFIGURATION

```
foo = 4
cpucore[0:31] = 'r10' # iacoma3
[r10]
fetchWidth = 4
issueWidth = $(foo)
branchDelay = $(foo)+2 # 6 minimum (deps enforced)
maxBranches = 4 # outstanding branches limit
maxIRequests= 2
iALUUnits = 2
iFPUnits = 2
iBJUnits = 1
LDSTQueue = 16 # Unified LD/ST buffer queue
winsize = 32
bb4Cycle = 1 # max BB fetch per cycle
bpred = 'r10kBPred'

<../benchs/branchModels.conf> # Include different branch models
```





# REPORT FORMAT

- By default is sesc\_bench.xxxxxxx
- Includes the input configuration file
- Homogeneous way to report data (GStats.h)

KEY:field1=value1:field2=value2:...

- **Example:**

```
ProcessId(0):totalTime=358100288:waitTime=0:spawnTime=0
Proc(0):clockTicks=358100289
BPred_RAS(0):nHit=      13
BPred_RAS(0):nMiss=     0
cache(0):readHit=      6788685
```



# REPORT OUTPUT

- Simulating an R10K at 180Mhz
- Crafty native execution 109msec

```
# Bench : ./mtst1 -h0x800000 ../benchs/crafty
# File  : sesc_crafty.EMomhe      :      Mon Mar 25 21:27:33 2002
      Sim. Speed      Exe Time      Sim Time (180MHz)
      1137.96 KIPS      17.88 secs      105.50 msec
Proc  Type      Total      RAS      BPred      BTB
   0  2bit      80.819% ( 99.91% of 10.05%) 84.02% 77.40%
File  IPC      Cycles      Busy      Control      Window      Struct      Memory
   0  1.07  18990598  26.79%  29.29%  18.45%  15.42%  0.20%  9.85%
```



# PROFILING

Function	nCalls	call(clk)	IPC	Acum(%)	Time(%)	Ctrl(%)
SCVQScores_al	998	1161690.1	1.19	23.34%	23.34%	25.79%
chan_v_eval	1568385	578.9	0.43	41.62%	18.28%	72.29%
_doscan	903173	581.6	1.56	52.20%	10.58%	68.07%
skipto	22804	9285.8	1.05	56.46%	4.26%	78.42%
hmm_tied_read	47	4469093.4	1.83	60.69%	4.23%	62.13%
root_chan_v_	386348	488.0	0.60	64.48%	3.80%	63.28%
.....						



# SIMULATION ACCURACY I

Simulating an R10K at 180Mhz (SGI Origin 200 IP27):

Bench	Native	Simulated	Error	SimTime
matrix	78ms	80ms	+2%	18s
crafty	109ms	105ms	-3%	18s
mcf	791ms	792ms	+0%	142s
mp3dec	643ms	619ms	-5%	160s
lat_mem	2405ms	2532ms	+5%	230s



# SIMULATION ACCURACY II

Simulating an R4400 at 150Mhz (SGI IP22):

Bench	Native	Simulated	Error	SimTime
matrix	291ms	282ms	-4%	27s
crafty	265ms	271ms	+4%	43s
mcf	2438ms	2422ms	-1%	174s
mp3dec	2347ms	2521ms	+7%	185s
lat_mem	4005ms	4129ms	+3%	279s





# SIMULATION SPEED

- Crafty simulation in a Pentium-III 1GHz

**1137KIPS** SESC simulating an R10K (with backend)



# SIMULATION SPEED

- Crafty simulation in a Pentium-III 1GHz

**1137KIPS** SESC simulating an R10K (with backend)

**178KIPS** smt simulating an R10K (no backend)





# SIMULATION SPEED

- Crafty simulation in a Pentium-III 1GHz

**1137KIPS** SESC simulating an R10K (with backend)

**178KIPS** smt simulating an R10K (no backend)

**2608KIPS** mint in-order





# SIMULATION SPEED

- Crafty simulation in a Pentium-III 1GHz

**1137KIPS** SESC simulating an R10K (with backend)

**178KIPS** smt simulating an R10K (no backend)

**2608KIPS** mint in-order

**≈150KIPS** simplescalr sim-outorder





# FULL SPECINT



- Assuming 1000KIPS we can run:
- 175.vpr Ref in  $\approx 30h$





# FULL SPECINT



- Assuming 1000KIPS we can run:
- 175.vpr Ref in  $\approx 30h$
- 197.parser Ref in  $\approx 90h$





# FULL SPECINT



- Assuming 1000KIPS we can run:
- 175.vpr Ref in  $\approx 30h$
- 197.parser Ref in  $\approx 90h$
- Reduced Data sets (Minnesota) in  $\approx 30min$



# EVEN FASTER

- Parallelization:
  - Low  $\approx 2$ : 2 threads likely to be balanced
  - High  $\approx 4$ : Several threads in multiprocessors



# EVEN FASTER

- Parallelization:
  - Low  $\approx 2$ : 2 threads likely to be balanced
  - High  $\approx 4$ : Several threads in multiprocessors
- Configuration Value Profile  $\approx 20\% - 40\%$ 
  - Convert variables to constants
  - The compiler simplifies the code
  - Run-time verification

Close to 3MIPS in our dual athlon



# EVEN FASTER

- Parallelization:
  - Low  $\approx 2$ : 2 threads likely to be balanced
  - High  $\approx 4$ : Several threads in multiprocessors
- Configuration Value Profile  $\approx 20\% - 40\%$ 
  - Convert variables to constants
  - The compiler simplifies the code
  - Run-time verification

Close to 3MIPS in our dual athlon



# DO YOU WANNA RUN IT?

- Add <http://sesc.sourceforge.net> to your bookmark
- Subscribe to [sesc-news@lists.sourceforge.net](mailto:sesc-news@lists.sourceforge.net)
- Read all the files in doc directory
- See the ./src/main.cpp example
- Understand the callbacks (DInst.h is easy)
- Understand the configuration files (\*.conf)





- More advance branch predictor: Multi-level, multi-cycle

- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)

- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)

# TODO

- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them



- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them
- Type a hackers guide (started)

- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them
- Type a hackers guide (started)
- Energy model?



- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them
- Type a hackers guide (started)
- Energy model?
- Give more flexibility to profiling

- More advance branch predictor: Multi-level, multi-cycle
- Add I-Cache (soon)
- Document the code with doxygen (starting)
- Create sample codes and document them
- Type a hackers guide (started)
- Energy model?
- Give more flexibility to profiling
- Use gcc instead of MIPSPro to generate benchmarks

